

Data visualization

COSC 480B

Reyan Ahmed

rahmed1@colgate.edu

Lecture 22


A peek into autoencoders

Overview

- Getting to know neural networks
- Designing autoencoders
- Representing images by using an autoencoder

Neural networks

A graphical representation of the linear equation $f(x) = w \times x + b$. The nodes are represented as circles, and edges are represented as arrows. The values on the edges are often called weights, and they act as a multiplication on the input. When two arrows lead to the same node, they act as a summation of the inputs.



Neural networks

Exercise 1

Is $f(x) = |x|$ a linear function?

Neural networks

Exercise 1


Is $f(x) = |x|$ a linear function?

ANSWER

No. It's two linear functions stitched together at zero, and that's not a single straight line.


Neural networks

Use nonlinear functions such as sig, tan, and ReLU to introduce nonlinearity to your models.




Neural networks

A nonlinear function, such as sigmoid, is applied to the output of a node.




Neural networks

A two-input network will have three parameters (w_1 , w_2 , and b). Remember, multiple lines leading to the same node indicate summation.




Neural networks

The input dimension can be arbitrarily long. For example, each pixel in a grayscale image can have a corresponding input x_i . This neural network uses all inputs to generate a single output number, which you might use for regression or classification. The notation w^T means you're transposing w , which is an $n \times 1$ vector, into a $1 \times n$ vector. That way, you can properly multiply it with x (which has the dimensions $n \times 1$). Such a matrix multiplication is also called a dot product, and it yields a scalar (one-dimensional) value.




Neural networks

Nodes that don't interface to both the input and the output are called hidden neurons. A hidden layer is a collection of hidden units that aren't connected to each other.




Autoencoders

If you want to create a network where the input equals the output, you can connect the corresponding nodes and set each parameter's weight to 1.



Autoencoders

Here, you introduce a restriction to a network that tries to reconstruct its input. Data will pass through a narrow channel, as illustrated by the hidden layer. In this example, there's only one node in the hidden layer. This network is trying to encode (and decode) an n -dimensional input signal into just one dimension, which will likely be difficult in practice.



Autoencoders

Exercise 2

Let x denote the input vector (x_1, x_2, \dots, x_n) , and let y denote the output vector (y_1, y_2, \dots, y_n) . Lastly, let w and w' denote the encoder and decoder weights, respectively. What's a possible cost function to train this neural network?

Autoencoders

Exercise 2

Let x denote the input vector (x_1, x_2, \dots, x_n) , and let y denote the output vector (y_1, y_2, \dots, y_n) . Lastly, let w and w' denote the encoder and decoder weights, respectively. What's a possible cost function to train this neural network?

ANSWER

See the loss function in listing 7.3.

Autoencoders

Python class schema

```
class Autoencoder:  
    def __init__(self, input_dim, hidden_dim):    1  
  
    def train(self, data):                        2  
  
    def test(self, data):                        3
```

- 1 Initializes variables
- 2 Trains on a dataset
- 3 Tests on some new data

Autoencoders

Using name scopes

```
with tf.name_scope('encode'):  
    weights = tf.Variable(tf.random_normal([input_dim,  
hidden_dim],  
dtype=tf.float32), name='weights')  
    biases = tf.Variable(tf.zeros([hidden_dim]), name='biases')
```

Autoencoders

Autoencoder class

```
import tensorflow as tf
import numpy as np

class Autoencoder:
    def __init__(self, input_dim, hidden_dim, epoch=250,
                 learning_rate=0.001):
        self.epoch = epoch
        self.learning_rate = learning_rate

        x = tf.placeholder(dtype=tf.float32, shape=[None,
input_dim])
```

- 1 Number of learning cycles
- 2 Hyperparameter of the optimizer
- 3 Defines the input layer dataset

Autoencoders

Autoencoder class

```
with tf.name_scope('encode'): 4
    weights = tf.Variable(tf.random_normal([input_dim,
hidden_dim],
    dtype=tf.float32), name='weights')
    biases = tf.Variable(tf.zeros([hidden_dim]),
name='biases')
    encoded = tf.nn.tanh(tf.matmul(x, weights) + biases)
    with tf.name_scope('decode'): 5
        weights = tf.Variable(tf.random_normal([hidden_dim,
input_dim],
        dtype=tf.float32), name='weights')
        biases = tf.Variable(tf.zeros([input_dim]),
name='biases')
        decoded = tf.matmul(encoded, weights) + biases
```

4 Defines the weights and biases under a name scope so you can tell them apart from the decoder's weights and biases

5 The decoder's weights and biases are defined under this name scope.

Autoencoders

Autoencoder class

```
self.x = x                                6
self.encoded = encoded                    6
self.decoded = decoded                    6

self.loss =
tf.sqrt(tf.reduce_mean(tf.square(tf.subtract(self.x,
self.decoded))))                          7
self.train_op =

tf.train.RMSPropOptimizer(self.learning_rate).minimize(self.l
oss)  8
self.saver = tf.train.Saver()              9
```

6 These will be method variables.

7 Defines the reconstruction cost

8 Chooses the optimizer

9 Sets up a saver to save model parameters as they're being learned

Autoencoders

Training the autoencoder

```
def train(self, data):
    num_samples = len(data)
    with tf.Session() as sess:
        sess.run(tf.global_variables_initializer())
        for i in range(self.epoch):
            for j in range(num_samples):
                l, _ = sess.run([self.loss, self.train_op],
                                feed_dict={self.x: [data[j]]})
                if i % 10 == 0:
                    print('epoch {0}: loss = {1}'.format(i, l))
                    self.saver.save(sess, './model.ckpt')
    self.saver.save(sess, './model.ckpt')
```

1 Starts a TensorFlow session, and initializes all variables

2 Iterates through the number of cycles defined in the constructor

3 One sample at a time, trains the neural network on a data item

4 Prints the reconstruction error once every 10 cycles

5 Saves the learned parameters to file

Autoencoders

Testing the model on data

```
def test(self, data):  
    with tf.Session() as sess:  
        self.saver.restore(sess, './model.ckpt')  
1  
        hidden, reconstructed = sess.run([self.encoded,  
self.decoded],  
        feed_dict={self.x: data})  
        print('input', data)  
        print('compressed', hidden)  
        print('reconstructed', reconstructed)  
        return reconstructed
```

- 1 Loads the learned parameters
- 2 Reconstructs the input

Autoencoders

Using your Autoencoder class

```
from autoencoder import Autoencoder
from sklearn import datasets

hidden_dim = 1
data = datasets.load_iris().data
input_dim = len(data[0])
ae = Autoencoder(input_dim, hidden_dim)
ae.train(data)
ae.test([[8, 4, 6, 2]])
```

Autoencoders

The test function shows info about the encoding and decoding process:

```
('input', [[8, 4, 6, 2]])  
('compressed', array([[ 0.78238308]], dtype=float32))  
('reconstructed', array([[ 6.87756062,  2.79838109,  
6.25144577,  
2.23120356]], dtype=float32))
```


Batch training

Batch helper function

```
def get_batch(X, size):  
    a = np.random.choice(len(X), size, replace=False)  
    return X[a]
```

Batch training

Batch learning


```
def train(self, data, batch_size=10):  
    with tf.Session() as sess:  
        sess.run(tf.global_variables_initializer())  
        for i in range(self.epoch):  
            for j in range(500):  
                batch_data = get_batch(data, self.batch_size)  
                l, _ = sess.run([self.loss, self.train_op],  
                                feed_dict={self.x: batch_data})  
                if i % 10 == 0:  
                    print('epoch {0}: loss = {1}'.format(i, l))  
                    self.saver.save(sess, './model.ckpt')  
            self.saver.save(sess, './model.ckpt')
```

1 Loops through various batch selections

2 Runs the optimizer on a randomly selected batch

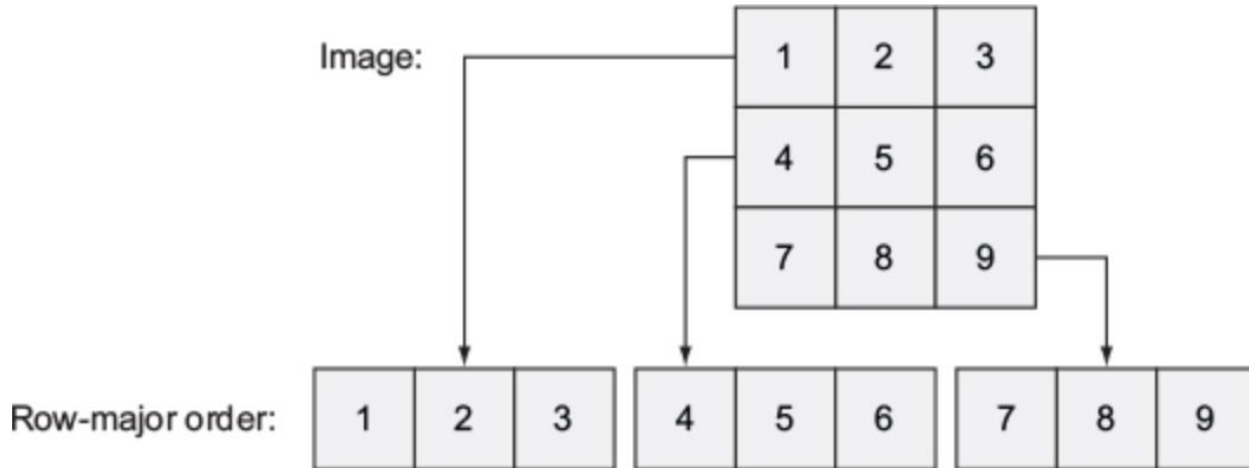
Working with images

A colored image is composed of pixels, and each pixel contains values for red, green, and blue.



Working with images

An image can be represented in row-major order. That way, you can represent a two-dimensional structure as a one dimensional structure.



Working with images

Loading images

```
from scipy.misc import imread, imresize

gray_image = imread(filepath, True)
small_gray_image = imresize(gray_image, 1. / 8.)
x = small_gray_image.flatten()
```

- 1 Loads an image as grayscale
- 2 Resizes it to something smaller
- 3 Converts it to a one-dimensional structure

Working with images

Exercise 3

Can you name other online image datasets? Search online and look around for more!

Working with images

Exercise 3

Can you name other online image datasets? Search online and look around for more!

ANSWER

Perhaps the most used in the deep-learning community is ImageNet (www.image-net.org). A great list can also be found online at <http://deeplearning.net/datasets>.

Working with images

Reading from the extracted CIFAR-10 dataset

```
import pickle

def unpickle(file):
    fo = open(file, 'rb')
    dict = pickle.load(fo, encoding='latin1')
    fo.close()
    return dict
```

1 Reads the CIFAR-10 file, returning the loaded dictionary

Working with images

Reading all CIFAR-10 files to memory

```
import numpy as np

names =
unpickle('./cifar-10-batches-py/batches.meta')['label_names']
data, labels = [], []
for i in range(1, 6):
    filename = './cifar-10-batches-py/data_batch_' + str(i)
    batch_data = unpickle(filename)
    if len(data) > 0:
        data = np.vstack((data, batch_data['data']))
        labels = np.hstack((labels, batch_data['labels']))
    else:
        data = batch_data['data']
        labels = batch_data['labels']
```

- 1 Loops through the six files
- 2 Loads the file to obtain a Python dictionary
- 3 The rows of a data sample represent each sample, so you stack it vertically.
- 4 Labels are one-dimensional, so you stack them horizontally.

Working with images

Converting CIFAR-10 image to grayscale

```
def grayscale(a):  
    return a.reshape(a.shape[0], 3, 32, 32).mean(1).reshape(a.shape[0], -1)  
  
data = grayscale(data)
```

Working with images

Setting up the autoencoder

```
from autoencoder import Autoencoder

x = np.matrix(data)
y = np.array(labels)

horse_indices = np.where(y == 7)[0]

horse_x = x[horse_indices]

print(np.shape(horse_x)) # (5000, 3072)

input_dim = np.shape(horse_x)[1]
hidden_dim = 100
ae = Autoencoder(input_dim, hidden_dim)
ae.train(horse_x)
```

Working with images

The output will trace loss values of every 10 epochs:

```
epoch 0: loss = 99.8635025024
epoch 10: loss = 35.3869667053
epoch 20: loss = 15.9411172867
epoch 30: loss = 7.66391372681
epoch 40: loss = 1.39575612545
epoch 50: loss = 0.00389165547676
epoch 60: loss = 0.00203850422986
epoch 70: loss = 0.00186171964742
epoch 80: loss = 0.00231492402963
epoch 90: loss = 0.00166488380637
epoch 100: loss = 0.00172081717756
epoch 110: loss = 0.0018497039564
epoch 120: loss = 0.00220602494664
epoch 130: loss = 0.00179589167237
epoch 140: loss = 0.00122790911701
epoch 150: loss = 0.0027100709267
epoch 160: loss = 0.00213225837797
epoch 170: loss = 0.00215123943053
epoch 180: loss = 0.00148373935372
epoch 190: loss = 0.00171591725666
```

Application of autoencoders

- A stacked autoencoder
- A denoising autoencoder
- A variational autoencoder